

Deploy Golang Two Tier Go Application on Kubernetes with CI/CD Pipeline

Author: Nitesh Nepali

DevOps Engineer/ Cyber Security Enthusiast: (CCNA, AWS-SAA, Linux (RHCSA), DevSecOps)

Email: pingnitesh07@gmail.com

Project Overview

Assessment:

Set up a CI/CD pipeline for a **Go** OR Java **Spring Boot** application using any CI/CD tools and containerize the application. The application must be deployed in a VM (you may use EC2 or GCP VM instance) using Kubernetes and helm. Implement monitoring to monitor your resource consumption.

Requirements:

1. **The entire system must be built using Terraform for dev and prod environments.**
2. Optimize resource utilization within the Kubernetes cluster. **Implement Horizontal Pod Autoscaling (HPA) based on CPU or custom metrics.**
3. The docker image must be optimized.
4. The application must be highly available and secure and optimize cost as much as possible.
5. Manage container replicas, and ensure high availability and scalability. **Configure Persistent Volume (PV) and Persistent Volume Claims (PVC).**
6. Configure health checks.

Some boilerplates for reference:

```
wget https://autostrada.dev/download/7u7gfr6nxhbai.zip
```

```
unzip 7u7gfr6nxhbai.zip
```

```
https://github.com/spring-projects/spring-boot
```

Tools and Services Used

This project "**Two Tier Go Application deploy on Kubernetes cluster** " involves a complete cloud-native setup using various AWS services, DevOps tools, and third-party platforms. Below is a categorized list of the tools and services used:

AWS Services

1. **Amazon VPC** – For custom network configuration with public and private subnets
2. **Amazon EC2** – To launch Ubuntu instances and host intermediate setup and for Jenkins
3. **IAM (Identity and Access Management)** – For creating users and assigning permissions
6. **Auto Scaling Group (ASG)** – To scale EC2 instances automatically
7. **Elastic Load Balancer (ELB)** – To distribute traffic and to Kubernetes cluster
8. **Amazon Route 53** – For domain name system (DNS) configuration

Infrastructure as a Code tool: [terraform](#) and [terraform vault](#)

Kubernetes cluster setup: [Kops](#) and [terraform](#)

Kubernetes package manager: [helm](#)

Containerization tool: [docker](#) and [docker hub](#)

Container orchestration: [Kubernetes](#)

CI/CD tool: [Jenkins](#) and [ArgoCD](#)

Monitoring and Alerting: [Prometheus](#) and [Grafana](#)

Code Repository and Version control System: [Git](#) and [GitHub](#)

Local Development & Testing Tools

1. **Git** – For cloning and version control
2. **Docker** – To containerize the application
3. **AWS CLI** – To interact with AWS services from the terminal

Third-Party Tools

1. **Mercantile domain registration** – For domain registration and DNS management
2. **Web-based Name Server Checker** – To verify Route 53 propagation

Phase 1: Initial Setup and Local Development

Step 1: Launch EC2 (Ubuntu 22.04):

- Provision an EC2 instance on AWS with Ubuntu 22.04.
- Connect to the instance using SSH.

Step 2: Clone the Code:

- Update all the packages and then clone the code.
- Clone your application's code repository onto the EC2 instance.

```
sudo apt update && upgrade -y
```

```
sudo apt install git -y
```

```
sudo git clone https://github.com/NiteshOne/Golang_codebase.git
```

Step 3: Install Golang, PostGreSql and Setup Database

A. Install golang and PostGreSql

- `sudo install golang -y`
- `sudo apt install postgresql postgresql-contrib -y`

B. Setup Database:

- a. Switch to the postgres user:
 - `sudo -i -u postgres`
- b. Access the PostgreSQL Shell
 - `psql`
- c. Create a New User (Role)
 - `CREATE USER myuser WITH PASSWORD 'mypassword';`
- d. Create a New Database
 - `CREATE DATABASE mydb OWNER myuser;`
- e. Create a New User (Role)
 - `ALTER DATABASE mydb OWNER TO myuser;`
- f. Grant Privileges to the User
 - `GRANT ALL PRIVILEGES ON DATABASE mydb TO myuser;`
- g. Exit PostgreSQL and postgres User.
 - `\q`
 - Exit

c. Export DB String as environment variable:

- `export DB_DSN="postgres://myuser:mypassword@localhost:5432/db?sslmode=disable"`

D. Build Golang code and run the app:

a. Navigate to the clone repo and download the dependencies

- `cd Golang_codebase`
- `go mod tidy`

b. Navigate to folder locating main.go

- `cd ~/Golang_codebase/cmd/web`
- `go run .`

E. Test the Connection on port 4444

`http://<ubuntu -IP>:<4444>` you will see the page shown in below



Phase 2: Infrastructure Setup

In this phase, Kubernetes cluster for dev and prod is provision using Kops and Terraform.

KOPS Configuration Procedures

Step 1: Launch EC2 (Ubuntu 22.04) as Management server:

- Provision an EC2 instance on AWS with Ubuntu 22.04.
- Connect to the instance using SSH.
- Management server is used for managing Kubernetes cluster.

Step 2: Install Kubectl

- `curl -LO https://storage.googleapis.com/kubernetes-release/release/$(curl -s https://storage.googleapis.com/kubernetes-release/release/stable.txt)/bin/linux/amd64/kubectl`
- `curl -LO https://storage.googleapis.com/kubernetes-release/release/v1.9.0/bin/linux/amd64/kubectl`
- `chmod +x ./kubectl`
- `sudo mv ./kubectl /usr/local/bin/kubectl`
- `export PATH=$PATH:/usr/local/bin/`
- `kubectl`

Note: run all command as root user

Step 2 - Install kops

- `#Install "kops" on Linux`
- `yum install wget`
- `curl -LO https://github.com/kubernetes/kops/releases/download/$(curl -s https://api.github.com/repos/kubernetes/kops/releases/latest | grep tag_name | cut -d '"' -f 4)/kops-linux-amd64`
- `chmod +x kops-linux-amd64`
- `sudo mv kops-linux-amd64 /usr/local/bin/kops`

- kops upgrade
Reference - <https://github.com/kubernetes/kops>

Step 3 - Configure AWS Route 53

DNS Changes (Refer the images show below;)

- Create a public Hosted Zone in Route53 k8s.niteshnepali.com.np
- Get a NS & SOA
- Go to Your domain provider and update the Name Server.
- Enter the NS records

The screenshot shows the AWS Route 53 console interface. On the left is a navigation menu with options like Route 53, Dashboard, Hosted zones, Health checks, Profiles, IP-based routing, Traffic flow, Domains, Resolver, DNS Firewall, and Application Recovery Controller. The main panel displays 'Hosted zone details' for the zone 'k8s.niteshnepali.com.np'. Below this, the 'Records (10)' tab is active, showing a table of DNS records. The table has columns for Record name, Type, Routing policy, Differ..., Alias, Value/Route traffic to, TTL (s...), and Health ...

Record name	Type	Routing policy	Differ...	Alias	Value/Route traffic to	TTL (s...)	Health ...
<input type="checkbox"/> k8s.niteshnepali.com.np	NS	Simple	-	No	ns-2016.awsdns-60.co.uk. ns-1331.awsdns-38.org. ns-804.awsdns-36.net. ns-434.awsdns-54.com.	172800	-
<input type="checkbox"/> k8s.niteshnepali.com.np	SOA	Simple	-	No	ns-2016.awsdns-60.co.uk. a...	900	-
<input type="checkbox"/> api.k8s.niteshnepali.com.np	A	Simple	-	Yes	api-k8s-niteshnepali-com--b...	-	-
<input type="checkbox"/> api.k8s.niteshnepali.com.np	AAAA	Simple	-	Yes	api-k8s-niteshnepali-com--b...	-	-
<input type="checkbox"/> argocd.k8s.niteshnepali.com.np	CNAME	Simple	-	No	a4fe3d12fba57432da80738f...	300	-
<input type="checkbox"/> bastion.k8s.niteshnepali.com.np	A	Simple	-	Yes	bastion-k8s-niteshnepali--trr...	-	-
<input type="checkbox"/> bastion.k8s.niteshnepali.com.np	AAAA	Simple	-	Yes	bastion-k8s-niteshnepali--trr...	-	-
<input type="checkbox"/> api.internal.k8s.niteshnepali.com.np	A	Simple	-	No	172.20.125.70	60	-
<input type="checkbox"/> kops-controller.internal.k8s.niteshnepali.com.np	A	Simple	-	No	172.20.125.70	60	-
<input type="checkbox"/> www.k8s.niteshnepali.com.np	A	Simple	-	Yes	dualstack.a4fe3d12fba5743...	-	-

Step 4 - Create a IAM Role or IAM User and Assign access to management Server

Create a IAM user or IAM Role with following Permission.

- AmazonEC2FullAccess
- AmazonRoute53FullAccess
- AmazonS3FullAccess
- IAMFullAccess
- AmazonVPCFullAccess
- AmazonSQSFullAccess

— AmazonEventBridgeFullAccess

In Case of IAM Role attach to the management server

In Case of IAM user configure access key and secret key

Step 6 - Craete a S3 bucket

Create and then list a new S3 bucket

```
$ aws s3 mb s3://k8s.niteshnepali.com.np
```

```
$ aws s3 ls
```

Step 7 – Generating a terraform configuration using Kops

```
kops create cluster \  
--name=k8s.niteshnepali.com.np \  
--state=s3://k8s.niteshnepali.com.np \  
--zones=ap-south-1a, ap-south-1b \  
--node-count=2 \  
--control-plane-count=1 \  
--node-size=t3.medium \  
--control-plane-size=t3.medium \  
--control-plane-zones=us-east-1a \  
--control-plane-volume-size=30 \  
--node-volume-size=30 \  
--topology=private \  
--networking=calico \  
--bastion \  
--ssh-public-key=~/.ssh/id_rsa.pub \ #ssh public key path  
--dns-zone=k8s.niteshnepali.com.np \  
--target=terraform \  
--out=terraform \  
--yes
```

Note: for remote access of the bastion host generate the ssh public key and provide the path

Step 8 – Creating Cluster

First install the terraform using official documentation.

- `cd terraform`
- `terraform init`
- `terraform plan`
- `terraform apply`

Step 9 – Updating the Cluster

- `kops get cluster --name k8s.niteshnepali.com.np -o yaml > cluster.yaml`
- `kops replace -f cluster.yaml --state=s3://niteshnepali.com.np`
- `kops update cluster --name k8s.niteshnepali.com.np --state=s3://k8.niteshnepali.com.np --out=./terraform --target=terraform`
then repeat Step no 8

The screenshot shows the AWS Management Console's EC2 Instances page. The left sidebar contains a navigation menu with categories: EC2, Images, Elastic Block Store, and Network & Security. The main content area is titled 'Instances (4) Info' and shows a table of four instances. All instances are in a 'Running' state. The table columns are: Name, Instance ID, Instance state, Instance type, Status check, Alarm status, and Availability. Below the table, there is a 'Select an instance' section.

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability
nodes-ap-south-1b.k8s.niteshnepali.com.np	i-057babb855a758afe	Running	t3.medium	3/3 checks passed	View alarms +	ap-south-1
nodes-ap-south-1a.k8s.niteshnepali.com.np	i-088a3f86cfa646e4	Running	t3.medium	3/3 checks passed	View alarms +	ap-south-1
control-plane-ap-south-1a.masters.k8s.niteshnepali.com.np	i-02a07edd7cb102372	Running	t3.medium	3/3 checks passed	View alarms +	ap-south-1
bastions.k8s.niteshnepali.com.np	i-0ab1e5a201c2b8e81	Running	t3.micro	3/3 checks passed	View alarms +	ap-south-1

Phase 3: Dockerizing the app and setup CI/CD

Dockerizing the golang code

Step 1. Install docker on management server

```
sudo apt install docker.io
```

Step 2. Clone the Repo

```
git clone https://github.com/NiteshOne/Golang\_codebase.git
```


Step 3. Create a DockerFile

```
1  # Build Stage
2  FROM golang as builder
3  #set environment variables
4  ENV CGO_ENABLED=0 GOOS=linux
5
6  WORKDIR /opt
7
8  # Copy source files
9  COPY . .
10
11 #download dependencies
12 RUN go mod tidy
13
14 # Build the Go application
15 RUN go build -v -o app ./cmd/web
16
17 #Final Stage
18 FROM debian:bullseye-slim
19
20 # Create non-root user and group
21 RUN groupadd -r goapp && useradd -r -g goapp goapp
22
23 WORKDIR /app
24
25 COPY --from=builder /opt/app /usr/local/bin/app
26
27 # Set permissions
28 RUN chown goapp:goapp /usr/local/bin/app
29
30 # Switch to non-root user
31 USER goapp
32
33 # Expose the port the app listens on
34 EXPOSE 4444
35
36 # Run the app
37 CMD ["app"]
38
```

Step 4. Install Jenkins and setup pipeline

- `sudo wget -O /etc/apt/keyrings/jenkins-keyring.asc \`
`https://pkg.jenkins.io/debian-stable/jenkins.io-2023.key`
`echo "deb [signed-by=/etc/apt/keyrings/jenkins-keyring.asc]" \`
`https://pkg.jenkins.io/debian-stable binary/ | sudo tee \`
`/etc/apt/sources.list.d/jenkins.list > /dev/null`
- `sudo apt-get update`
- `sudo apt-get install Jenkins`

Note: For the Jenkins installation and java compatible follow Jenkins official documentation

Step 5. Create a Jenkins file and run the pipeline

First, create a docker hub credentials and configure on Jenkins to use it

```
1 pipeline{
2   agent {label 'linux' }
3
4   tools{
5     git 'Default'
6   }
7
8   environment {
9     IMAGE_NAME = 'nitace/go-lang-app'
10    IMAGE_TAG = "v${env.BUILD_NUMBER}"
11    REGISTRY_CREDENTIALS = 'docker-hub-credentials'
12  }
13  stages{
14
15    stage ('Build Image'){
16      steps{
17        sh 'docker build -t ${IMAGE_NAME}:${IMAGE_TAG} .'
18      }
19    }
20
21    stage ('PUSH TO REGISTRY'){
22      steps {
23        withCredentials([usernamePassword(credentialsId: "${env.REGISTRY_CREDENTIALS}", usernameVariable: 'USERNAME', passwordVariable: 'PASSWORD')]) {
24          script {
25            sh """
26              echo "$PASSWORD" | docker login -u "$USERNAME" --password-stdin
27              docker push ${IMAGE_NAME}:${IMAGE_TAG}
28            """
29          }
30        }
31      }
32    }
33  }
34 }
```

Phase 3: Monitoring (Prometheus and Grafana)

We are using helm package manager to install Prometheus and Grafana.

Step 1. Install helm on the Management Server

- `$ curl -fsSL -o get_helm.sh https://raw.githubusercontent.com/helm/helm/main/scripts/get-helm-3`
- `$ chmod 700 get_helm.sh`
- `$./get_helm.sh`

Step 2. Add Prometheus Repo and update it and installed it

- `helm repo add prometheus-community https://prometheus-community.github.io/helm-charts`
- `helm repo update`
- `helm install prometheus prometheus-community/kube-prometheus-stack --namespace monitoring --create-namespace`

Get Grafana 'admin' user password by running:

```
kubectl --namespace monitoring get secrets prometheus-grafana -o jsonpath="{.data.admin-password}" | base64 -d ; echo
```

Step 3. Exposing service via ingress

- `helm repo add ingress-nginx https://kubernetes.github.io/ingress-nginx`
- `helm repo update`
- `helm install ingress-nginx ingress-nginx/ingress-nginx --namespace ingress-nginx --create-namespace`

Step 4. Create ingress file for both Prometheus and Grafana

For Grafana

create grafana_ingress.yaml

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: grafana-ingress
  namespace: monitoring
  annotations:
    kubernetes.io/ingress.className: "nginx" # or "alb" if using AWS ALB ingressI
    nginx.ingress.kubernetes.io/rewrite-target: /
    # Optional: enable basic auth or TLS here
spec:
  rules:
  - host: grafana.niteshnepali.com.np
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: prometheus-grafana
            port:
              number: 80
```

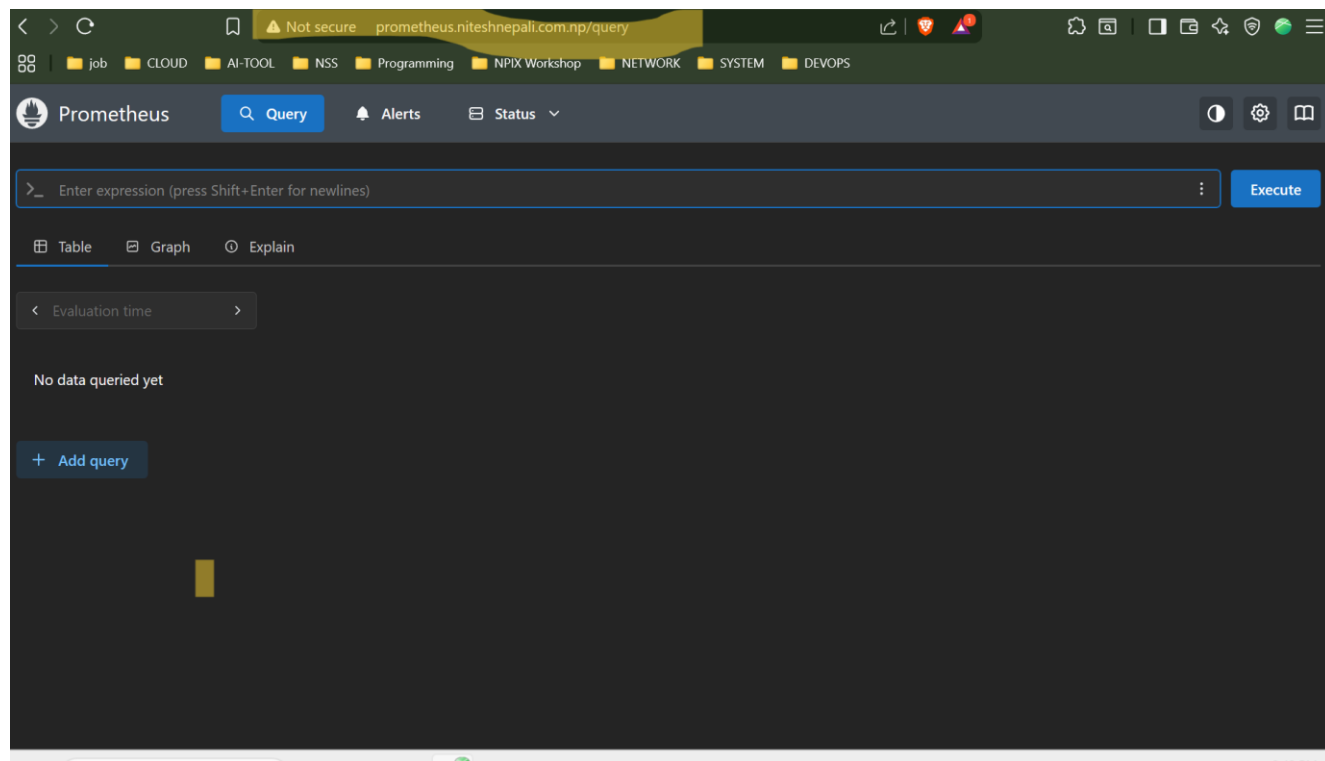
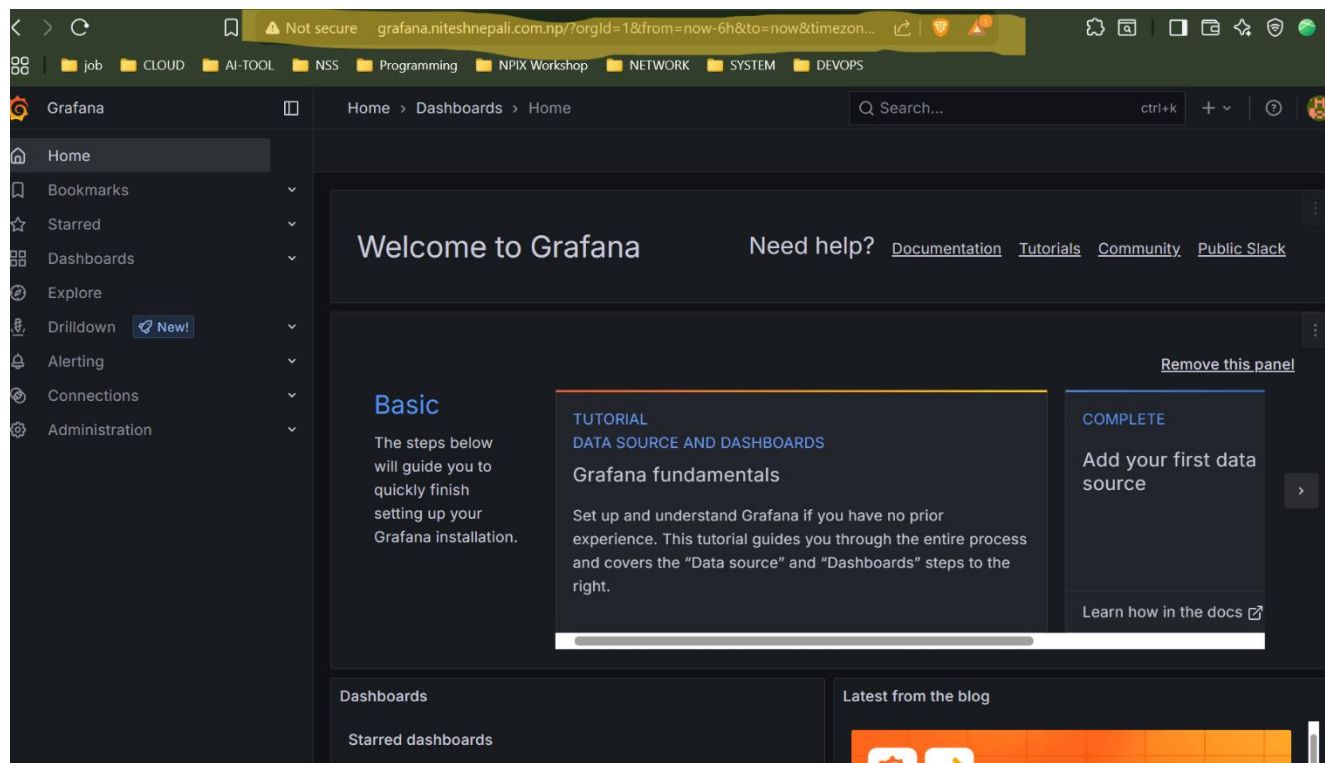
For Prometheus

Create prometheues_ingress.yaml

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: prometheus-ingress
  namespace: monitoring
  annotations: {}
spec:
  ingressClassName: nginx
  rules:
  - host: prometheus.niteshnepali.com.np
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: prometheus-kube-prometheus-prometheus
            port:
              number: 9090
```

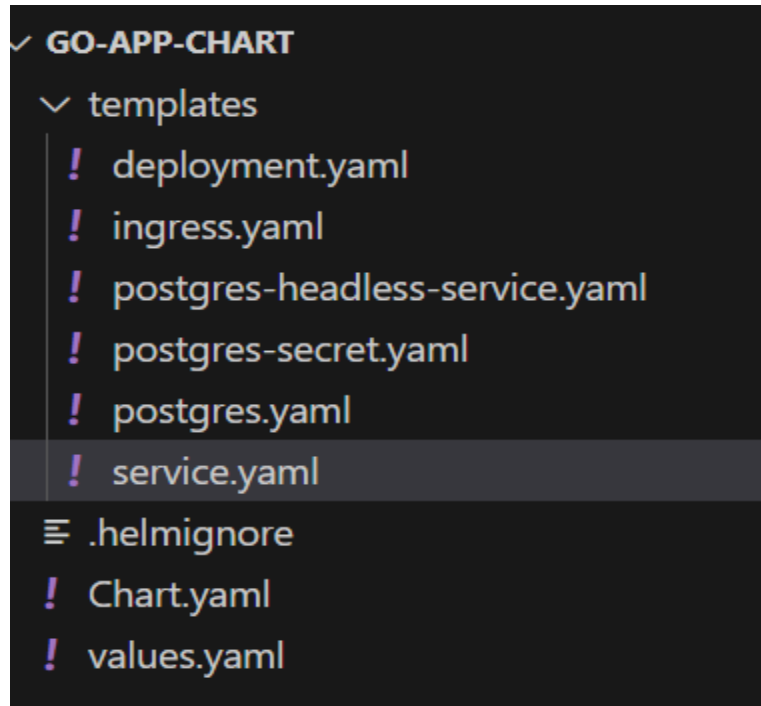
Step 5. Create a CNAME record in Route53 and point to ingress controller and access it.

<input type="checkbox"/>	niteshnepali.com.np	SOA	Simple	-	No	ns-1303.awsdns-34.org.
<input type="checkbox"/>	grafana.niteshnepali.com.np	CNAME	Simple	-	No	ns-1692.awsdns-19.co.uk. awsdns-hostmaster.amazon.cor
<input type="checkbox"/>	k8s.niteshnepali.com.np	NS	Simple	-	No	ns-2016.awsdns-60.co.uk. ns-1331.awsdns-38.org. ns-804.awsdns-36.net. ns-434.awsdns-54.com.
<input type="checkbox"/>	prometheus.niteshnepali.com.np	CNAME	Simple	-	No	a4fe3d12fba57432da80738fab32c5dd-1078502515.ap-s



Phase 5. Creating helm chart for Kubernetes Deployment

Creating a helm chart Folder Structure



template/deployment.yaml

```
# templates/deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: {{ .Release.Name }}-deploy
  labels:
    env: production
    tier: backend
spec:
  replicas: {{ .Values.replicaCount }}
  selector:
    matchLabels:
      env: production
      tier: backend
  template:
    metadata:
      labels:
        env: production
        tier: backend
    spec:
      initContainers:
        - name: wait-for-postgres
          image: busybox
          command:
            - sh
            - -c
            - |
              until nc -z postgres-0.postgres-headless.default.svc.cluster.local 5432; do
                echo "Waiting for Postgres TCP connection..."
                sleep 3
              done
      dnsPolicy: ClusterFirst
      containers:
        - name: testcont
          image: "{{ .Values.image.repository }}:{{ .Values.image.tag }}"
          ports:
            - containerPort: 4444
          env:
            - name: DB_DSN
              valueFrom:
                secretKeyRef:
                  name: postgres-secret
                  key: DB_DSN
          readinessProbe:
            tcpSocket:
              port: 4444
            initialDelaySeconds: 5
            periodSeconds: 10
            failureThreshold: 3
          livenessProbe:
            tcpSocket:
              port: 4444
            initialDelaySeconds: 15
            periodSeconds: 20
            failureThreshold: 3
```


template/service.yaml

```
# templates/service.yaml
apiVersion: v1
kind: Service
metadata:
  name: {{ .Release.Name }}-svc
  labels:
    env: production
    tier: backend
spec:
  type: {{ .Values.service.type }}
  selector:
    env: production
    tier: backend
  ports:
    - protocol: TCP
      port: {{ .Values.service.port }}
      targetPort: {{ .Values.service.targetPort }}
```

Template/postgres.yaml

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: {{ .Values.postgres.name }}
  labels:
    app: {{ .Values.postgres.name }}
spec:
  serviceName: "{{ .Values.postgres.serviceName }}"
  replicas: {{ .Values.postgres.replicaCount }}
  selector:
    matchLabels:
      app: {{ .Values.postgres.name }}
  template:
    metadata:
      labels:
        app: {{ .Values.postgres.name }}
    spec:
      containers:
        - name: {{ .Values.postgres.name }}
          image: "{{ .Values.postgres.image }}"
          imagePullPolicy: IfNotPresent
          envFrom:
            - secretRef:
                name: {{ .Values.postgres.secretName }}
          ports:
            - containerPort: 5432
          volumeMounts:
            - name: postgres-data
              mountPath: /var/lib/postgresql/data3
      volumeClaimTemplates:
        - metadata:
            name: postgres-data
          spec:
            accessModes: ["ReadWriteOnce"]
            resources:
              requests:
                storage: {{ .Values.postgres.storage }}
```

postgres-headless-service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: {{ .Values.postgres.serviceName }}
spec:
  clusterIP: None
  selector:
    app: {{ .Values.postgres.name }}
  ports:
    - port: 5432
      name: db
```

postgres-secret.yaml

```
apiVersion: v1
kind: Secret
metadata:
  name: {{ .Values.postgres.secretName }}
type: Opaque
stringData:
  POSTGRES_USER: {{ .Values.postgres.auth.user | quote }}
  POSTGRES_PASSWORD: {{ .Values.postgres.auth.password |
quote }}
  POSTGRES_DB: {{ .Values.postgres.auth.database | quote }}
  DB_DSN: "user={{ .Values.postgres.auth.user }} password={{
.Values.postgres.auth.password }} host={{
.Values.postgres.dsnHost }} port=5432 dbname={{
.Values.postgres.auth.database }} sslmode=disable"
```

Ingress.yaml

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: {{ .Release.Name }}-ingress
  labels:
    app: {{ .Release.Name }}
  annotations:
    kubernetes.io/ingress.class: nginx
spec:
  ingressClassName: nginx
  rules:
    - host: {{ .Values.ingress.host | default "example.com"
    }}
    http:
      paths:
        - path: /
          pathType: Prefix
          backend:
            service:
              name: {{ .Release.Name }}-svc
              port:
                number: {{ .Values.service.port | default
80 }}
```

Phase 6: Setup ArgoCD and deployment

Step 1. Create Namespace for ArgoCD

- `kubectl create namespace argocd`

Step 2. Add ArgoCD Helm Repository

- `helm repo add argo https://argoproj.github.io/argo-helm`
- `helm repo update`

Step 3. Install ArgoCD via Helm

- `helm install argocd argo/argo-cd \`
`--namespace argocd`

Step 4. Expose argocd via ingress

Argocd-ingress.yaml

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: argocd-ingress
  namespace: argocd
  annotations:
    nginx.ingress.kubernetes.io/backend-protocol: "HTTPS"
spec:
  ingressClassName: nginx
  rules:
  - host: argocd.k8s.niteshnepali.com.np
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: argocd-server
            port:
              number: 443
```

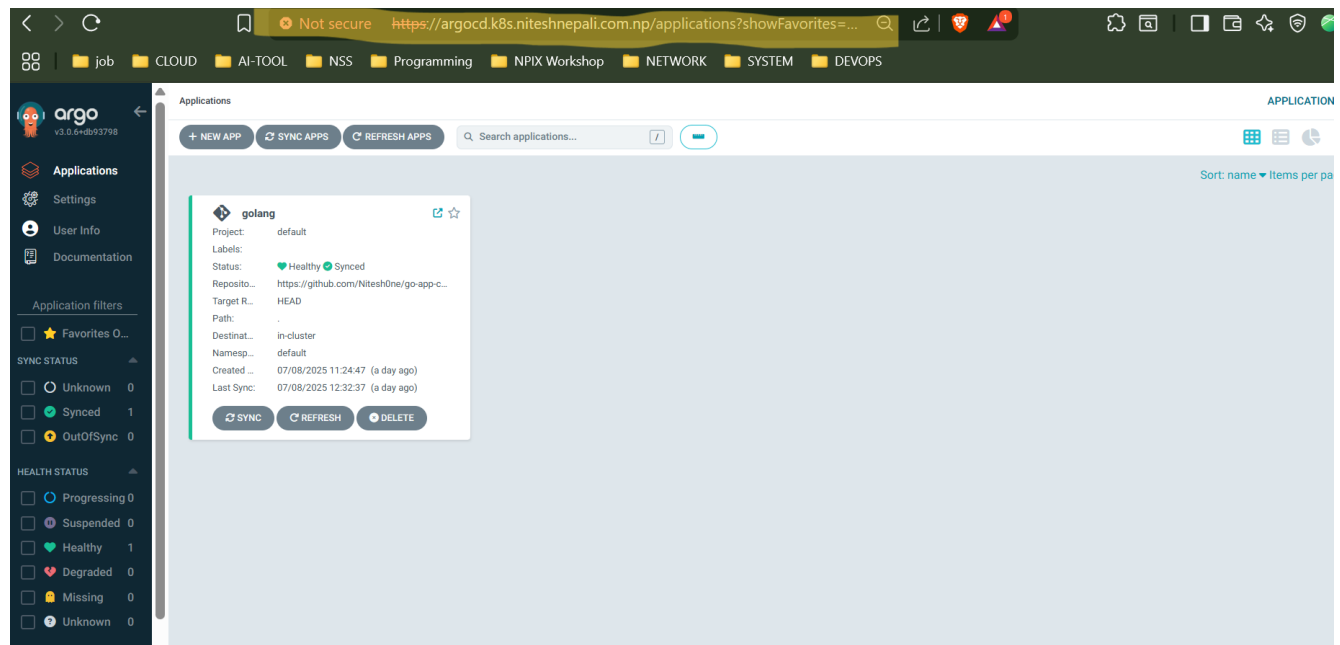
Step 5. Create Route53 CNAME Record and point to ingress controller

<input type="checkbox"/>	api.k8s.niteshnepali.com.np	AAAA	Simple	-	Yes	api-k8s-niteshnepali-com--b503ee-4e03a8b6c0d0f619.elb.ap-sout...
<input type="checkbox"/>	argocd.k8s.niteshnepali.com.np	CNAME	Simple	-	No	a4fe3d12fba57432da80738fab32c5dd-1078502515.ap-south-1.elb...

Step 6: Access argocd

Get argocd credentials running command

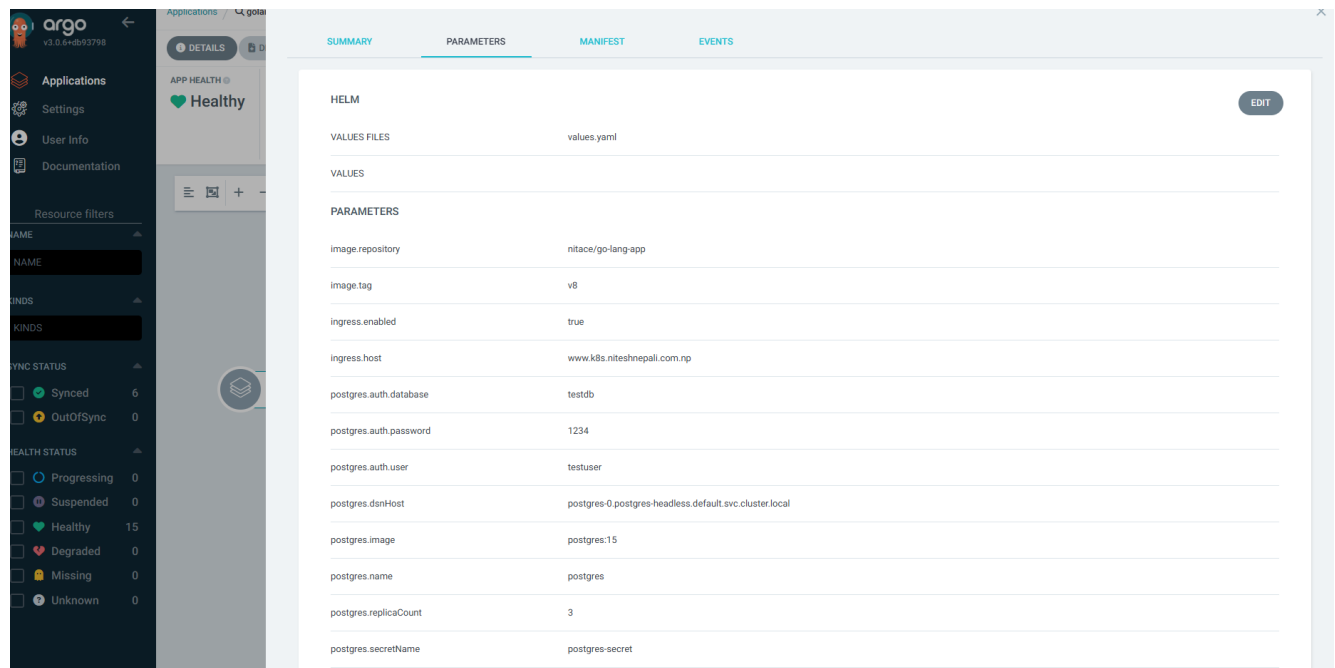
```
kubectl -n argocd get secret argocd-initial-admin-secret -o jsonpath="{.data.password}" | base64 -d
```



Phase 7: Deployment using argocd

Step 1. Create a helm repo and push to the github

Step 2. Confiure the argocd for deployment

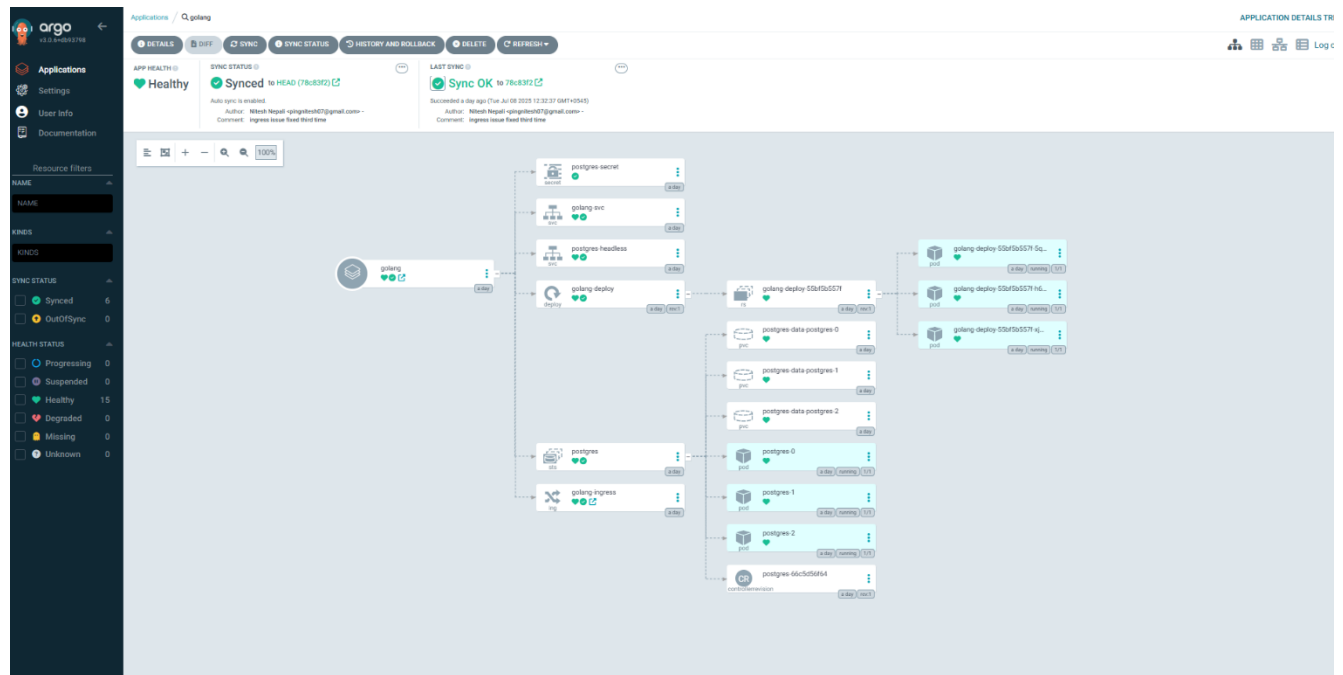


Step 6. Verify the deployment

3 pod of the golang container

3 pod of the postgres database

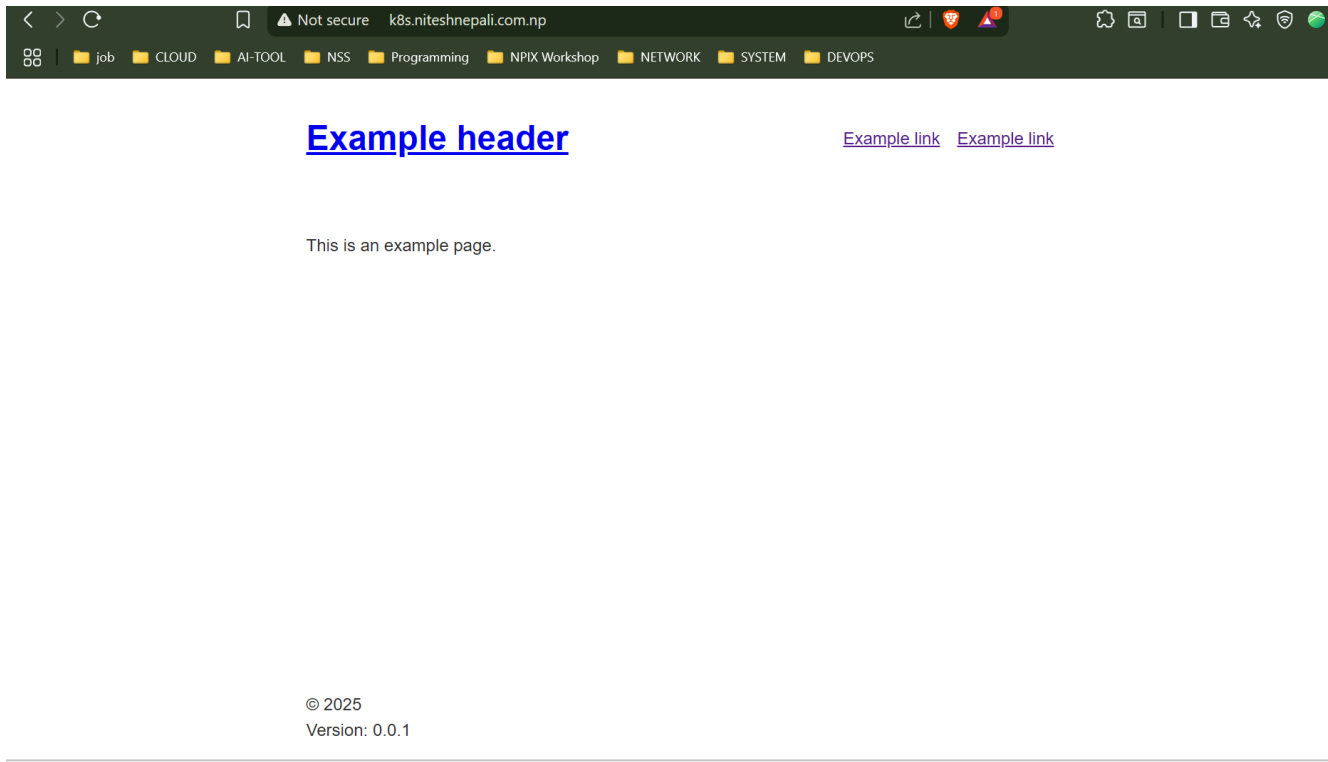
Ingress controller for the application,



Step 6: create the Route53 CNAME record for the application and access it

www.k8s.niteshnepali.com.np	A	Simple	-	Yes	dualstack.a4fe3d12fba57432da80738fab32c5dd-1078502515.ap-s...
-----------------------------	---	--------	---	-----	---

Step7. Access the application



Project Repo:

Kubernetes cluster using kops and Terraform repo:

https://github.com/NiteshOne/kops_cluster_infra_setup.git

Go Application with Dockerfile and Jenkins File:

https://github.com/NiteshOne/Golang_codebase.git

Kubernetes helm chart repo:

<https://github.com/NiteshOne/go-app-chart.git>